

# BarCat Deals: Detailed Design (D1)

**Team Members:** Kaustubh Mathur, Brian Nguyen, Ved Sanap, Meredith Bartel, Kiki Vasilev

**Course:** Senior Design

**Advisor:** Dr. Hrishikesh Vinayak Bhide

**Assignment:** Design D1 - Detailed Design

**Date:** October 7, 2026

---

## 1. Header, Scope, and Conventions

### Project Title

**BarCat Deals**

### Goal Statement

BarCat Deals gives University of Cincinnati students an easier way to find deals around campus. By using AI, the system automatically finds deals in businesses' social media posts, and business owners can also submit their own deals. Students can search deals that are active right now and save the ones they like.

**Basic input:** Social media posts processed by an AI, and manual deal submissions from business owners.

**Basic output:** Deals from bars and restaurants around campus, served to students as JSON.

### Scope

This document details the **Backend API Service (C3)** and the **Deal Ingestion Service (C5)**, together with the data they use in the **Deal Database (C4)**. It covers the deal extraction and de-duplication algorithm and the active-deal search algorithm. The Student Web App (C1), Owner Web Portal (C2), and External Social Media APIs (C6) are deferred to D2.

### Diagram Conventions

- Rectangular boxes represent database tables (entities) and their attributes. Only attributes that matter for behavior are shown.
- PK identifies a primary key. FK identifies a foreign key that references a record in another table. UK identifies a unique constraint.

- Lines between tables show relationships. Crow's foot symbols show cardinality: || means exactly one, o| means zero or one, and o{ means zero or more.
- Times use local Eastern time ("HH:MM") unless a field is marked UTC.
- D0 component names (C1 to C6) and interface IDs (I1 to I5) keep the meaning they have in D0.

## 2. Data Model, the D1 Diagram

### 2.1 Entity-Relationship Diagram

BarCat Deals stores every deal in a shared base table, plus one extra table for each source: owner submissions and automated extractions. Raw social media posts are stored separately, so one post can produce several deals and the original text is kept for debugging.



### 2.2 Entity Descriptions

| Entity            | Purpose                                                                                                          |
|-------------------|------------------------------------------------------------------------------------------------------------------|
| STUDENT           | A student user. <code>age_boolean</code> is true once the student confirms they are 21 or older.                 |
| SAVED_DEAL        | A bookmark linking one student to one deal.                                                                      |
| BUSINESS_OWNER    | An owner who manages one venue and may submit deals for it.                                                      |
| VENUE             | A bar or restaurant near campus, with its address and social media account ID.                                   |
| SOCIAL_MEDIA_POST | A raw post fetched from a venue's account, with its processing status.                                           |
| BASE_DEAL         | The common record for every deal: description, category, weekdays, start and end time, alcohol flag, and status. |
| SUBMITTED_DEAL    | Extra data for a deal an owner entered by hand: which owner submitted it.                                        |
| AUTOMATED_DEAL    | Extra data for a deal found by the AI: its source post, confidence score, and last scraped time.                 |

## 2.3 Structural Decisions

| Decision                                                     | Rationale                                                                                                                                   |
|--------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------|
| Relational database, not key-value or document               | Every entity links to others by keys, and everyone on the team knows SQL.                                                                   |
| VENUE is its own entity                                      | A venue has its own attributes and hosts many deals and posts over time. Keeping it separate avoids repeating values on every deal row.     |
| SOCIAL_MEDIA_POST is its own entity                          | Raw posts are stored before parsing. This keeps history for debugging, and one post can contain several deals.                              |
| SAVED_DEAL is its own entity                                 | A student saves many deals and a deal is saved by many students, so an explicit join table is needed.                                       |
| BASE_DEAL split into SUBMITTED_DEAL and AUTOMATED_DEAL       | Both kinds share description, category, times, and the alcohol flag, but each has its own source data. Splitting avoids many empty columns. |
| Weekdays, start time, and end time stored as separate fields | "Is this deal active now?" becomes a simple comparison and needs no text parsing.                                                           |
| venue_id on BUSINESS_OWNER                                   | The API can check that an owner submits deals only for their own venue.                                                                     |

| Decision                         | Rationale                                                                                                  |
|----------------------------------|------------------------------------------------------------------------------------------------------------|
| status on <code>BASE_DEAL</code> | Low-confidence automated deals are stored but never shown, and old deals are hidden without being deleted. |

## 2.4 Indexing Decisions

BarCat Deals indexes `BASE_DEAL.alcohol`, because every student search filters on it. Every foreign key (`venue_id`, `student_id`, `base_deal_id`, `owner_id`, `post_id`) is indexed because listing a student's saved deals and a venue's deals needs constant joins.

`AUTOMATED_DEAL.scraped_at` is indexed so old deals can be found and hidden quickly. The unique pair (`source_platform`, `external_id`) on `SOCIAL_MEDIA_POST` stops the same post from being stored twice. Free-text fields are not indexed because nothing searches them, and every extra index slows writes.

## 2.5 Deal Lifecycle

1. C5 fetches a post and stores it with status `new`.
2. C5 extracts deals from the post and sends them to C3.
3. C3 sets each deal's status from its confidence score: `published`, `pending_review`, or `rejected`.
4. Owner submissions skip steps 1 to 3 and are `published` immediately.
5. A background job sets a deal to `expired` after its `valid_until` date.
6. Only `published` deals are returned by search.

---

## 3. Core Algorithms

Two computations break the product if they are wrong. Extraction fills the feed with wrong, duplicate, or unflagged alcohol deals if it fails. Search shows students deals that are over, or alcohol deals they should not see, if it fails. Neither is heavy work at our size, so the risk is correctness, not speed.

**Assumed data size:** about 50 venues posting about 2 times a day gives about 100 posts a day. About 300 deals are live at a time, and about 50,000 deal rows exist after a year. At 100 times that, there are about 10,000 posts a day and about 5 million rows.

### 3.1 Algorithm A: Deal Extraction and De-duplication

#### Purpose

Turns one social media post into zero or more structured deals without storing the same deal twice. C5 extracts and checks each deal. C3 saves it, because only C3 talks to the database.

## Inputs and Outputs

| Parameter | Type          | Description                                |
|-----------|---------------|--------------------------------------------|
| post_id   | string        | A post in <code>SOCIAL_MEDIA_POST</code> . |
| venue_id  | string        | The venue that posted.                     |
| text      | string        | Post caption, 0 to 5,000 characters.       |
| images    | array of URLs | 0 to 10 flyer images.                      |
| timestamp | UTC datetime  | Used to turn "tonight" into a real date.   |

Output:

- An array of 0 to 10 deals. Each has `title` (string, up to 120 characters), `deal_description` (string, up to 1,000), `category` (`food`, `drink`, or `other`), `alcohol` (boolean), `days_of_week` (integers 1 to 7), `start_time` and `end_time` ("HH:MM"), `valid_until` (date or null), and `confidence_score` (float from 0 to 1).
- One result per deal: `inserted`, `merged`, `pending_review`, or `rejected`. The post becomes `processed`, `no_deal`, or `failed`.

## Approach

1. C5 picks up posts with status `new`, one worker per post. A post with no text and no images becomes `no_deal`.
2. C5 sends the text, images, and post date to the AI model and asks for JSON with explicit dates.
3. C5 checks the JSON: end time differs from start time, dates are real and within 60 days, and at least one weekday is set.
4. `alcohol` is true if the AI says so or a keyword matches (beer, wells, shots, draft, happy hour, 21+). When unsure, it is true.
5. C3 applies a confidence gate: 0.70 or higher is `published`, 0.40 to 0.69 is `pending_review`, and below 0.40 is `rejected`. These thresholds are starting values to tune.
6. C3 compares the deal with the same venue's live deals. If weekdays and times match and the title words are at least 80% alike (Jaccard similarity), it merges: `scraped_at` is updated and the higher confidence is kept. Otherwise it inserts a new deal.

## Complexity

Each post needs one AI call (about 2 to 10 seconds), so about 100 calls a day. Merging compares against about 20 live deals for one venue, so it takes  $O(k)$  time for  $(k)$  candidates. At 100 times the volume, that is about 7 posts a minute, which a few workers can handle. Only cost and rate limits grow, so the algorithm does not need tuning.

## Why This Approach

Regex or rule-based parsing breaks on emoji, flyers, and varied wording. A trained model would need labeled posts we do not have. Exact-match de-duplication misses reworded re-posts, and embeddings are too heavy for about 20 candidates. An AI model with a fixed JSON format needs no training data and costs little at this volume, and Jaccard matching is simple and easy to test.

## Edge Cases

- **Empty post or no deal:** Mark `no_deal` and store nothing.
- **Several deals in one post:** Save each one with the same `post_id`.
- **Same post fetched twice:** Blocked by the unique pair `(source_platform, external_id)`.
- **Same deal re-posted or posted on two platforms:** Merge it.
- **Two existing deals both match:** Merge into the most recently scraped one.
- **Unclear dates or "until close":** Lower the confidence so the deal goes to review.
- **Bad AI output:** Retry once. After 3 attempts, mark the post `failed`.
- **AI service down:** Leave the post `new` and retry later with a growing delay. Never publish unchecked output.

## 3.2 Algorithm B: Active-Deal Search with the 21+ Gate

### Purpose

Returns the deals that are active right now, filtered and sorted, and never returns alcohol deals to a student who has not confirmed they are 21 or older. This runs on every page load.

### Inputs and Outputs

| Parameter                                                       | Type                                                                      | Description                         |
|-----------------------------------------------------------------|---------------------------------------------------------------------------|-------------------------------------|
| <code>at</code>                                                 | <code>datetime with offset</code>                                         | Time to search at. Defaults to now. |
| <code>venue_id</code> ,<br><code>category</code>                | <code>string</code> , <code>enum</code>                                   | Optional filters.                   |
| <code>include_alcohol</code>                                    | <code>boolean</code>                                                      | Defaults to false.                  |
| <code>sort</code> , <code>limit</code> ,<br><code>cursor</code> | <code>enum</code> , <code>integer 1 to 50</code> ,<br><code>string</code> | Ordering and paging.                |

| Parameter | Type                 | Description                                         |
|-----------|----------------------|-----------------------------------------------------|
| caller    | anonymous or student | Taken from the login token, never from the request. |

Output:

- An array of deals, each with `minutes_remaining` (integer, minutes), and a `next_cursor` (string or null).

## Approach

1. Convert `at` to Eastern time to get a weekday and a time of day.
2. Keep deals that are `published`, not past `valid_until`, and whose weekday and time window contain now. For overnight deals (end earlier than start), also match when the previous weekday is set and the time is before the end.
3. Include alcohol deals only if `include_alcohol` is true and the caller is a logged-in student with `age_boolean` true.
4. Sort food deals first, then by `minutes_remaining`, then by `id`. Page with a cursor, not offsets.

## Complexity

About 300 live deals, so one indexed query takes about a millisecond. At 100 times (about 30,000 live deals) it is still fast, because indexes on `status` and `alcohol` keep the scan to live rows. We will time the query with 100 times test data before tuning anything.

## Why This Approach

Creating one row per future occurrence would need a job to generate rows, and it would have to be redone on every edit. Text schedules such as cron strings need a parser we would have to maintain. Offset paging skips or repeats rows when deals expire between pages. Weekday flags with explicit start and end times need no parsing and no generator job.

## Edge Cases

- **No matching deals:** Return `200` with an empty list, not an error.
- **Overnight deals:** A Thursday 9 PM to 2 AM deal still shows at 1 AM Friday.
- **Daylight saving:** Handled by the standard time zone library using local clock time.
- **Anonymous caller or `age_boolean` false asks for alcohol:** Return `403`, so the app can prompt for confirmation.
- **Expired or future-dated deals:** Excluded, and a background job sets `status` to `expired`.

- **Bad at , limit , or cursor :** Return 422 .
- **Database timeout:** Return 503 .
- **Ties in the sort:** Broken by id .

## 4. Build-versus-Reuse Decisions

| Component                          | Build or Reuse | Library or Service             | License / Terms                           | Reason                                                                |
|------------------------------------|----------------|--------------------------------|-------------------------------------------|-----------------------------------------------------------------------|
| Relational database                | Reuse          | PostgreSQL                     | PostgreSQL License                        | Mature; supports the indexes and transactions we need.                |
| Table design, constraints, indexes | Build          | None                           | Project license to be selected            | Comes from our requirements.                                          |
| Database access and migrations     | Reuse          | SQLAlchemy and Alembic         | MIT                                       | Mature; parameterized queries block SQL injection.                    |
| API framework                      | Reuse          | FastAPI                        | MIT                                       | Mature and async; generates API docs automatically.                   |
| Request validation                 | Reuse          | Pydantic                       | MIT                                       | Checks types and ranges from the API contract.                        |
| Login and sessions                 | Reuse          | Supabase Auth (or Auth0)       | MIT (self-hostable); hosted-service terms | We do not hand-roll passwords. Free tier limits to be checked.        |
| Collecting posts (C5 from C6)      | Reuse          | Official Meta Graph and X APIs | Platform terms                            | Stable. Unofficial scrapers break platform terms and get blocked.     |
| Reading posts into deals           | Reuse          | Anthropic API and Python SDK   | SDK MIT; API terms                        | Handles messy captions and flyers; low cost at about 100 calls a day. |
| Dates and time zones               | Reuse          | zoneinfo , python-dateutil     | PSF; Apache and BSD                       | No hand-written date logic.                                           |

| Component                                       | Build or Reuse | Library or Service               | License / Terms                | Reason                                                  |
|-------------------------------------------------|----------------|----------------------------------|--------------------------------|---------------------------------------------------------|
| Scheduling                                      | Reuse          | APScheduler                      | MIT                            | Simple polling with no extra service.                   |
| HTTP calls                                      | Reuse          | httpx                            | BSD-3-Clause                   | Timeouts and retries built in.                          |
| Active-deal search, 21+ gate, owner-venue check | Build          | Project implementation in Python | Project license to be selected | Specific to our product and the part we must get right. |
| Duplicate matching                              | Build          | Python standard library          | Project license to be selected | About 20 lines, so a dependency is not worth it.        |
| Automated tests                                 | Reuse          | pytest                           | MIT                            | Fixtures suit seeding test data.                        |

The selected libraries and services should be checked for active maintenance, documentation, license compatibility, performance, and fit with the team's environment. License names above are from memory, so each teammate should open the project's repository and confirm the license file before final submission. Official social media APIs also limit reading other accounts' posts, so the owner submission path is the fallback. Confirm what each API allows before D2.

## 5. API Contract

C3 exposes endpoints to the clients (interfaces I1 and I2) and to C5 (interface I3). All bodies are JSON. Weekdays are 1 (Monday) to 7 (Sunday). Every error body has the shape `{"error": {"code": string, "message": string}}`.

### 5.1 Endpoints

| Method and Endpoint    | Inputs                                                                                                                | Success Response                                                                                               | Error Response                                                                                      |
|------------------------|-----------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------|
| GET /api/v1/deals (C1) | Optional query: at datetime with offset (default now, at most 14 days ahead); venue_id string; category (food, drink, | 200: data array (each with deal_id string, venue {venue_id, name, address}, title, deal_description, category, | 401 token invalid or expired<br>400 include_a from an anc caller or a s<br>age_boole<br>422 invalid |

| Method and Endpoint                                                           | Inputs                                                                                                                                                                                                                                                                                                                                                        | Success Response                                                                                                                                                                                                                                                      | Error Response                                                                                                                                      |
|-------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------|
|                                                                               | other );<br>include_alcohol<br>boolean (default<br>false); sort<br>( default ,<br>ending_soon ,<br>venue_name );<br>limit integer 1 to<br>50 (default 20);<br>cursor string.<br>Optional bearer<br>token                                                                                                                                                      | alcohol boolean,<br>source ( owner or<br>automated ),<br>days_of_week<br>integer array,<br>start_time ,<br>end_time ,<br>minutes_remaining<br>integer in minutes,<br>valid_until date<br>or null),<br>next_cursor string<br>or null,<br>generated_at UTC<br>timestamp | 429 too many requests;<br>503 database unavailable                                                                                                  |
| GET<br>/api/v1/deals/{deal_id}<br>(C1)                                        | deal_id string,<br>required                                                                                                                                                                                                                                                                                                                                   | 200 : one deal,<br>same fields as above<br>without<br>minutes_remaining                                                                                                                                                                                               | 403 alcohol not<br>caller not owner;<br>404 deal not<br>published;<br>database unavailable                                                          |
| POST /api/v1/owner/deals<br>(C2, owner login)                                 | title string 1 to<br>120;<br>deal_description<br>string 1 to 1,000;<br>category enum;<br>alcohol boolean;<br>days_of_week<br>array of 1 to 7<br>unique integers<br>from 1 to 7;<br>start_time ,<br>end_time<br>"HH:MM" (not<br>equal); optional<br>valid_until date,<br>at most 1 year<br>ahead. Venue<br>comes from the<br>owner's record,<br>never the body | 201 : the new deal<br>with status:<br>"published"                                                                                                                                                                                                                     | 400 malformed<br>401 not signed<br>caller is not<br>owner has<br>409 duplicate<br>(includes the<br>deal_id );<br>field; 503 database<br>unavailable |
| PUT and DELETE<br>/api/v1/me/saved-<br>deals/{deal_id} (C1, student<br>login) | deal_id string,<br>required                                                                                                                                                                                                                                                                                                                                   | PUT : 200 with<br>{"deal_id":<br>string, "saved":                                                                                                                                                                                                                     | 401 not signed<br>deal not found<br>database unavailable                                                                                            |

| Method and Endpoint                                                 | Inputs                                                                               | Success Response                                                                  | Error Response                                      |
|---------------------------------------------------------------------|--------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------|-----------------------------------------------------|
|                                                                     |                                                                                      | <pre> true}. DELETE: 204 </pre>                                                   |                                                     |
| POST<br>/api/v1/internal/automated-deals (C5, header X-Service-Key) | post_id string;<br>deals array of 0 to 10 extracted deals (fields as in Algorithm A) | 200: results array with index integer, disposition string, deal_id string or null | 401 invalid<br>404 post r<br>422 invalid database u |

Sending a post that is already processed replays the stored results and creates nothing new.

## 5.2 Example Request and Response

Example request:

```

GET /api/v1/deals?category=drink&include_alcohol=true&limit=1
Authorization: Bearer <token issued after sign-in>
Accept: application/json

```

Example success response:

```

{
  "data": [
    {
      "deal_id": "d_0412",
      "venue": { "venue_id": "v_07", "name": "Example Bar", "address": "123 Calhoun St" },
      "title": "$3 Wells",
      "deal_description": "$3 well drinks until close",
      "category": "drink",
      "alcohol": true,
      "source": "automated",
      "days_of_week": [3],
      "start_time": "21:00",
      "end_time": "02:00",
      "minutes_remaining": 275,
      "valid_until": null
    }
  ],
  "next_cursor": null,
  "generated_at": "2026-10-08T01:25:00Z"
}

```

The values are illustrative. `generated_at` is Wednesday 9:25 PM Eastern, so 275 minutes remain until 2:00 AM.

## 5.3 Age Confirmation Handling

Alcohol deals are filtered on the server using the stored `age_boolean`, never a value sent by the client. An anonymous caller or a student who has not confirmed they are 21 or older gets `403` when asking for alcohol deals, so the app can prompt for confirmation instead of silently hiding them. Deals are never filtered only in the browser.

## 5.4 Versioning

The API uses the `/api/v1` prefix. Adding an endpoint or an optional response field is not a breaking change. Removing or renaming a field, changing a field's type or unit, making an optional input required, or changing what an error code means is a breaking change and requires `/api/v2`, while `/v1` keeps working for a notice period the team sets.

---

# 6. Technology Choices with Justification

## 6.1 Database: PostgreSQL

PostgreSQL is the selected database because every BarCat Deals entity is related to others and we need unique constraints. **Team skill fit:** Brian Nguyen is the team's most experienced member with PostgreSQL, which he worked with during his co-op, and everyone on the team has used SQL and databases. **Licensing:** The PostgreSQL License is permissive and free. **Community support:** It is widely used, with full official documentation and many answered questions. **Performance:** Our data is small (about 50,000 rows now, about 5 million at 100 times) and the frequent queries are indexed lookups, so responses take milliseconds. **Cost and hosting:** The software is free, and Azure offers a managed PostgreSQL service (check current pricing and any student credits). MongoDB was considered as an alternative, but our data is relational and we need unique constraints and transactions.

## 6.2 Backend: Python with FastAPI

Python with FastAPI is the selected backend for C3 and C5, so both can share validation code. **Team skill fit:** Kiki Vasilev became proficient in Python through his experience at Siemens, and Kaustubh Mathur also writes Python (pandas, Tkinter automation). FastAPI is a thin layer over Python, so that skill carries over. **Licensing:** Python (PSF), FastAPI, Pydantic, and SQLAlchemy (MIT) have no fees or restrictions. **Community support:** All are widely used with strong documentation, and the Anthropic SDK supports Python. **Performance:** Peak load is a few requests a second, and the slow part is the AI call, which runs in the background. **Cost and**

**hosting:** It runs in one small container on Azure. Node.js with Express was considered, but we chose one language across C3 and C5. Django was also considered, but we do not need its admin site.

## 6.3 Front End: Flask

Flask is the selected framework for the student app (C1) and owner portal (C2), which D2 will detail. Flask renders the pages with Jinja templates and gets its data from C3 over HTTP. **Team skill fit:** Meredith Bartel has a lot of experience building websites, and Flask is written in Python, the same language as the backend, so teammates can help each other. **Licensing:** Flask uses the BSD-3-Clause license. **Community support:** Flask is mature, with extensive documentation and many examples. **Performance:** Server-rendered forms and lists are light and load fast on a phone. Flask's built-in server is for development only, so we will run it under a production server such as Gunicorn (MIT). **Cost and hosting:** It runs in a container on Azure next to the backend, with no separate front-end build or host. React with Vite was considered, but it adds a second language and a separate build toolchain that we do not need for forms and lists.

## 6.4 Job Processing: Database Table as the Queue

C5 uses a `status` column on `SOCIAL_MEDIA_POST` as its work queue instead of a separate message broker. Ved Sanap is the team member for this area. **Team skill fit:** The team already knows SQL, so this approach is easy to understand and debug. **Licensing:** Nothing extra to license. **Community support:** Picking rows with `FOR UPDATE SKIP LOCKED` is a well-known PostgreSQL pattern, and APScheduler (MIT) handles the polling schedule. **Performance:** 100 posts a day (10,000 at 100 times) is far below what this approach handles. **Cost and hosting:** No second service to pay for or host. Redis with Celery was considered, but a broker adds cost and moving parts we do not need.

## 6.5 Hosting: Microsoft Azure

The backend, the ingestion worker, and the Flask front end run in containers on Azure (for example Azure App Service or Azure Container Apps), connected to Azure's managed PostgreSQL service. **Team skill fit:** The whole team learned Azure in our Cloud Computing class, so nobody needs to learn a new platform. **Licensing:** Azure's service terms apply and our code stays ours. The Anthropic API is a commercial service. **Community support:** Azure has extensive official documentation and a large user community. **Performance:** One small instance covers campus traffic, and the main delay is the AI call, which is not on the student's request path. **Cost and hosting:** Cost is mostly Azure usage plus AI usage, which at about 100 short calls a day should be small. Check current pricing and any student credits, and track spending, because our Constraints essay lists API and hosting cost as a limit. Render and

Railway were considered, but we already know Azure from class. Raw virtual machines were also considered, but they add patching and security work for no benefit at this size.